

Supplementary Online Material for: Conditional modal branch selection in physics-informed active control of lid-driven cavity flow

Beniaiche Ahmed, Larabi Abderrahim, and Belkadi Mbarek
Ecole Militaire Polytechnique, Algiers, Algeria

This document provides the implementation details, complete datasets, and additional analyses supporting the main manuscript. Section numbers S1–S11 correspond to the references “Supplementary Sec. S...” used in the text.

1 S1. PINN architecture and training details

The parametric network represents the flow-control pair for the full Reynolds interval $[100, 1000]$ through inputs $(x, y, t, \widetilde{Re})$, with $\widetilde{Re} = (Re - 550)/450 \in [-1, 1]$.

- **Flow subnetwork:** 6 hidden layers $\times$ 64 neurons, tanh activation, 21 379 trainable parameters; outputs (u, v, p) .
- **Control subnetwork:** 4 hidden layers $\times$ 32 neurons, 3 614 trainable parameters; outputs the coefficients $\{a_{mn}\}$.
- **Precision:** all arithmetic in double precision; residuals by automatic differentiation.
- **Training:** 300-epoch pre-training (flow only, control frozen at zero), then 3 000-epoch joint phase with Adam ($\eta_0 = 8 \times 10^{-4}$, decay factor 0.995 per 100 epochs). Collocation points: $N_{\text{phys}} = 4\,000$ interior points resampled uniformly every epoch; $N_{\text{probe}} = 20$ for the Reynolds-separation term.
- **Loss weights:** $\lambda_{\text{PDE}} = 1$, $\lambda_{\text{BC}} = 15$, $\lambda_{\text{diss}} = 0.05$, $\lambda_{\text{ctrl}} = 200$, $\lambda_{\text{var}} = 1$, $\lambda_{Re} = 2$.

2 S2. Modal-energy quadrature conventions

For the lid-expansion $U_{\text{lid}}(x, t, Re) = \sum_m \sum_n a_{mn} \sin(m\pi x/L_x) \cos(n\pi t)$, the modal energy accounting uses $\langle \cos^2(0) \rangle = 1$ and $\langle \cos^2(n\pi t) \rangle = \frac{1}{2}$ for $n \geq 1$:

$$\mathcal{E}_{m0} = \frac{1}{2} a_{m0}^2, \quad \mathcal{E}_{mn} = \frac{1}{4} a_{mn}^2 \quad (n \geq 1).$$

The realized actuation energy is $E_r = \langle U_{\text{lid}}^2 \rangle_{s,t}$, with $s = x/L_x$, and the normalized modal fractions $f_{mn} = \mathcal{E}_{mn} / \sum \mathcal{E}_{ij}$ sum to unity. The stationary second-harmonic fraction is $f_2 \equiv f_{20}$ and the total time-dependent fraction is $f_t = \sum_m \sum_{n \geq 1} f_{mn}$.

Table 1: Pooled and per-condition statistics of the 185-realization campaign (30 cells; oversampled cells: $Re = 1000, E^* = 1$ ($n = 20$), $Re = 500, E^* = 0.25$ ($n = 15$), $Re = 700, E^* = 2$ ($n = 15$)).

Quantity	B_1	B_2	B_{first}	B_3
Pooled frequency (frozen labels)	42.2%	43.8%	0.0%	14.1%
Pooled frequency (threshold re-classification)	42.2%	43.8%	0.0%	14.1%
Mean per condition	42.1%	46.7%	0.0%	11.3%
Pooled excluding oversampled cells ($N = 135$)	42.2%	48.1%	0.0%	9.6%
Discrepancies frozen vs. threshold ($N = 185$)	0 / 185			

3 S3. Reference (Re, E^*) campaign – realization-level data

The threshold classification rule is: B_{first} if $f_1 > 0.5$; B_1 if $f_2 > 2f_t$ and $f_2 > 0.5$; B_2 if $f_t > 2f_2$ and $f_t > 0.5$; otherwise B_3 . No realization of the 185-run reference campaign at $\lambda_{\text{var}} = 1$ satisfies $f_1 > 0.5$, so B_{first} is empty there. The frozen labels and the threshold-based re-classification coincide for all 185 realizations (0/185). The realization-level file (Supplementary Table S1 in the submission system) lists, for each realization: (Re, E^*), seed, f_2 , f_t , the dominant coefficients, total loss, and frozen branch label.

4 S4. Variance-regularization campaign and Wilson intervals

Six values of $\lambda_{\text{var}} \in \{0, 0.25, 0.5, 1, 2, 5\}$ at three operating points, $N_s = 5$ seeds each (90 runs). Frequencies are reported in Table 3 of the main text. Finite-sample 95% Wilson-score confidence intervals for an estimated frequency $\hat{p} = k/N_s$ use the standard center-plus-half-width formulation

$$C = \frac{\hat{p} + z^2/(2N_s)}{1 + z^2/N_s}, \quad h = \frac{z}{1 + z^2/N_s} \sqrt{\frac{\hat{p}(1 - \hat{p})}{N_s} + \frac{z^2}{4N_s^2}}, \quad z = 1.96,$$

with interval $[\max(0, C-h), \min(1, C+h)]$. For $N_s = 5$ the resulting intervals are: 0/5: [0.000, 0.434], 1/5: [0.036, 0.625], 2/5: [0.118, 0.769], 3/5: [0.231, 0.882], 4/5: [0.375, 0.964], and 5/5: [0.566, 1.000].

All realizations are classified with the four-class rule of Sec. S3. At $\lambda_{\text{var}} = 0$ every one of the 15 realizations is B_{first} ($f_1 > 0.5$); these were grouped with B_3 under the original three-class scheme. At (700, 2) and $\lambda_{\text{var}} = 0.25$ the first-harmonic branch also survives in 5/5 realizations ($f_1 \in [0.544, 0.646]$), which were likewise labeled B_3 (mixed) under the original scheme and are identified as B_{first} under the corrected rule. Table S4b identifies every B_{first} realization of the 90-run campaign by its first-harmonic fraction and corrected label.

5 S5. Energy sweep and seed multiplicity

Energy sweep at $Re = 500$ ($E^* \in \{0.01, 0.05, 0.1, 0.25, 0.5, 1, 2\}$):

In every energy-sweep realization the first-harmonic fraction remains small ($f_1 \leq 8\%$), so the B_{first} branch never appears; the sweep stays within the B_1 branch and the erosion of second-harmonic purity at high energy is a redistribution toward time-dependent modes rather than toward the fundamental.

Seed multiplicity at $Re = 500$, $E^* = 0.25$, 10 seeds: 4 B_1 , 4 B_2 , 2 B_3 (0 B_{first} , $f_1 < 0.02$ for all seeds); total losses in [1.371, 1.394] (relative spread 1.7%).

Table 2: 90-run λ_{var} campaign: per-cell empirical frequency k/N_s and 95% Wilson-score interval for each of the four branch classes (corresponding to Table 3 of the main text).

(Re, E^*)	λ_{var}	B_1	B_2	B_{first}	B_3
(500,0.25)	0.0	0/5 —	0/5 —	5/5 [0.566, 1.000]	0/5 —
(500,0.25)	0.25	3/5 [0.231, 0.882]	2/5 [0.118, 0.769]	0/5 —	0/5 —
(500,0.25)	0.5	2/5 [0.118, 0.769]	1/5 [0.036, 0.625]	0/5 —	2/5 [0.118, 0.769]
(500,0.25)	1.0	2/5 [0.118, 0.769]	1/5 [0.036, 0.625]	0/5 —	2/5 [0.118, 0.769]
(500,0.25)	2.0	2/5 [0.118, 0.769]	2/5 [0.118, 0.769]	0/5 —	1/5 [0.036, 0.625]
(500,0.25)	5.0	2/5 [0.118, 0.769]	3/5 [0.231, 0.882]	0/5 —	0/5 —
(700,2.0)	0.0	0/5 —	0/5 —	5/5 [0.566, 1.000]	0/5 —
(700,2.0)	0.25	0/5 —	0/5 —	5/5 [0.566, 1.000]	0/5 —
(700,2.0)	0.5	0/5 —	3/5 [0.231, 0.882]	0/5 —	2/5 [0.118, 0.769]
(700,2.0)	1.0	1/5 [0.036, 0.625]	2/5 [0.118, 0.769]	0/5 —	2/5 [0.118, 0.769]
(700,2.0)	2.0	3/5 [0.231, 0.882]	1/5 [0.036, 0.625]	0/5 —	1/5 [0.036, 0.625]
(700,2.0)	5.0	3/5 [0.231, 0.882]	1/5 [0.036, 0.625]	0/5 —	1/5 [0.036, 0.625]
(1000,1.0)	0.0	0/5 —	0/5 —	5/5 [0.566, 1.000]	0/5 —
(1000,1.0)	0.25	0/5 —	3/5 [0.231, 0.882]	0/5 —	2/5 [0.118, 0.769]
(1000,1.0)	0.5	2/5 [0.118, 0.769]	2/5 [0.118, 0.769]	0/5 —	1/5 [0.036, 0.625]
(1000,1.0)	1.0	2/5 [0.118, 0.769]	2/5 [0.118, 0.769]	0/5 —	1/5 [0.036, 0.625]
(1000,1.0)	2.0	2/5 [0.118, 0.769]	2/5 [0.118, 0.769]	0/5 —	1/5 [0.036, 0.625]
(1000,1.0)	5.0	2/5 [0.118, 0.769]	3/5 [0.231, 0.882]	0/5 —	0/5 —

6 S6. Capacity and truncation experiments

Capacity variants at $Re = 500$, $E^* = 0.25$: 4×32 (small), 6×64 (reference), 8×128 (large) give $f_2 \in [89.5\%, 92.2\%]$. Control truncations $(N_x, N_t) \in \{(4, 3), (6, 1), (6, 3), (6, 5), (8, 5), (8, 9)\}$ give $f_2 \in [87.3\%, 92.3\%]$ with $A_2 \in [0.668, 0.684]$; see Table 6 of the main text.

7 S7. Spectral bases and aspect ratio

Fourier basis: sine in space (enforces corner compatibility) and cosine in time ($n = 0 = \text{stationary}$). Modified-Chebyshev basis: same structure with Chebyshev polynomials of the first kind T_m on the lid interval, normalized to the same energy budget. Aspect-ratio cells $AR \in \{0.5, 1, 2\}$ at $Re = 500$, $E^* = 0.25$: $f_2 = 69.9\%$, 92.2% , 91.2% (see Table 7 of the main text).

8 S8. MRT-LBM solver and grid-convergence index (GCI)

Solver. D2Q9 lattice with the diagonal relaxation vector $\mathbf{S} = \text{diag}(0, 1.4, 1.4, 0, 1.2, 0, 1.2, \omega, \omega)$, $\omega = 1/\tau$, $\tau = 3\nu_{\text{LBM}} + 0.5$, $\nu_{\text{LBM}} = U_{\text{ref}}N/Re$ with $U_{\text{ref}} = 0.05$. Lid forcing through a velocity boundary condition; boundary populations via the Zou–He scheme. Parallelism via Numba **prange** over lattice rows.

Benchmark. Centerline velocity errors against Ghia et al. at $N = 256$: $L_2(u) = 0.00506$, $L_\infty(u) = 0.01283$ ($Re = 100$); $L_2(u) = 0.01411$, $L_\infty(u) = 0.03003$ ($Re = 1000$); corresponding v -component errors $0.00675/0.01136$ and $0.01321/0.02125$.

GCI. Three-grid study of the $A_2 \sin(2\pi x)$ control at $Re = 500$, $N \in \{128, 256, 512\}$:

$$\varepsilon_{128} = 0.0093500, \quad \varepsilon_{256} = 0.0104070, \quad \varepsilon_{512} = 0.0109658.$$

Table 3: Per-realization data for the 20 B_{first} outcomes of the 90-run campaign: 15 realizations at $\lambda_{\text{var}} = 0$ and 5 at $(Re, E^*) = (700, 2)$ with $\lambda_{\text{var}} = 0.25$. f_1 , f_2 , and f_t are the first-harmonic, second-harmonic, and total time-dependent fractions.

Re	E^*	λ_{var}	seed	f_1	f_2	f_t
500	0.25	0.0	0	0.8988	0.0031	0.0541
500	0.25	0.0	1	0.9357	0.0002	0.0301
500	0.25	0.0	2	0.9107	0.0002	0.0467
500	0.25	0.0	3	0.9290	0.0007	0.0343
500	0.25	0.0	4	0.9251	0.0002	0.0399
700	2.0	0.0	0	0.9862	0.0011	0.0112
700	2.0	0.0	1	0.9817	0.0041	0.0135
700	2.0	0.0	2	0.9306	0.0241	0.0446
700	2.0	0.0	3	0.9539	0.0222	0.0193
700	2.0	0.0	4	0.9768	0.0020	0.0191
1000	1.0	0.0	0	0.9682	0.0062	0.0220
1000	1.0	0.0	1	0.9709	0.0102	0.0182
1000	1.0	0.0	2	0.8680	0.0012	0.1212
1000	1.0	0.0	3	0.9721	0.0073	0.0203
1000	1.0	0.0	4	0.9610	0.0005	0.0322
700	2.0	0.25	0	0.6456	0.0047	0.1712
700	2.0	0.25	1	0.5801	0.0988	0.2677
700	2.0	0.25	2	0.5443	0.1897	0.1766
700	2.0	0.25	3	0.6029	0.0517	0.3214
700	2.0	0.25	4	0.6397	0.0079	0.3123

Table 4: Energy sweep at $Re = 500$.

E^*	0.01	0.05	0.1	0.25	0.5	1	2
f_2 (%)	~ 94	~ 93	92.2	92.2	–	–	74.0
f_t (%)	~ 0.5	~ 0.7	1.1	1.3	–	–	21.8

Observed order $p_{\text{obs}} = 0.9196$ (from $p = \ln |(\varepsilon_{256} - \varepsilon_{128}) / (\varepsilon_{512} - \varepsilon_{256})| / \ln 2$), Richardson extrapolation $\varepsilon_{\infty} = 0.01159$. The GCI (formula S8.4) uses the safety factor $F_s = 1.25$ and refinement ratio $r = 2$:

$$\text{GCI}_j = F_s \frac{|\varepsilon_{j+1} - \varepsilon_j|}{(r^p - 1) \varepsilon_j}, \quad \Rightarrow \quad \text{GCI}_{512} = 7.15\%, \quad \text{GCI}_{256} = 14.25\%.$$

Uniform-to-extrapolated reduction: $0.01159 / 0.0442824 = 73.8\%$ (vs. 76.5% measured at $N = 256$).

9 S9. Symbolic regression: configuration and LOCO scores

PySR over binary operators $\{+, -, \times, /\}$ and unary $\{\sin, \cos, \exp, \text{square}\}$, 20 populations, 100 iterations each, Pareto selection on complexity–accuracy trade-off. Best expression: $U_{\text{SR}} \approx a[b + \cos(\omega_{\text{SR}}x - \phi)]^2(c_0 + c_1x + c_2t + c_3Re) + c_4$, $\omega_{\text{SR}} \approx 3.42$. Leave-one-Reynolds-condition-out R^2 over $Re \in \{100, 300, 400, 500, 600, 700, 800, 900, 1000\}$: minimum 0.9326 ($Re = 900$), maximum 0.9908 ($Re = 500$); all nine values above 0.93.

10 S10. Reproducibility double runs and multi-seed audit

Double runs A/B (idempotent, same seeds/hyperparameters, base 6×5 , $E^* = 0.25$, final code) at $Re = 500$:

Table 5: Reproducibility double-run results (A/B/original).

seed	run	E	f_2	f_t	A_{10}	A_{20}	branch
4	A	0.2514	0.5957	0.3693	0.0376	-0.5472	B_3
4	B	0.2512	0.6080	0.3575	0.0394	-0.5527	B_3
4	orig	0.2495	0.5215	0.4501	0.0364	-0.5101	B_3
7	A	0.2508	0.9271	0.0083	-0.0604	0.6819	B_1
7	B	0.2508	0.9271	0.0083	-0.0604	0.6819	B_1
7	orig	0.2498	0.9131	0.0162	-0.0681	0.6754	B_1
8	A	0.2507	0.0357	0.9609	-0.0109	0.1339	B_2
8	B	0.2507	0.0357	0.9609	-0.0109	0.1339	B_2
8	orig	0.2503	0.0133	0.9841	-0.0158	0.0817	B_2

Runs A and B are bit-identical to the last decimal for seeds 7 and 8, and quasi-identical for seed 4 (same class B_3). All three coincide with the original realizations re-analyzed with the final code.

Multi-seed audit (99 realizations, final code, 18 configurations):

Table 6: Audit summary per configuration: mean f_2 , mean f_t , mean A_2 (in %), and branch counts.

label	n	$\bar{f}_2$	$\bar{f}_t$	$\bar{A}_2$	$B_1/B_2/B_3$
ar_rect_2	5	90.8	1.2	0.687	5/0/0
ar_rect_half	5	15.3	83.8	0.187	1/4/0
ar_square	5	73.0	20.1	0.572	4/1/0
bas_cheb_re100	5	0.5	99.2	-0.030	0/5/0
bas_cheb_re1000	5	1.2	98.4	-0.047	0/5/0
bas_cheb_re500	5	2.1	97.6	-0.085	0/5/0
bas_fourier_re100	5	32.6	64.7	0.313	1/3/1
bas_fourier_re1000	5	73.5	21.3	0.610	3/0/2
bas_fourier_re500	5	73.0	20.1	0.572	4/1/0
cap_baseline	8	64.0	30.5	0.402	5/2/1
cap_large	8	89.1	3.6	0.676	8/0/0
cap_small	8	92.2	1.3	0.694	8/0/0
mc_4x3	5	54.4	39.4	0.436	3/2/0
mc_6x1	5	91.1	0.0	0.683	5/0/0
mc_6x3	5	75.1	16.1	0.485	4/1/0
mc_6x5	5	73.1	20.1	0.572	4/1/0
mc_8x5	5	59.1	34.9	0.485	3/2/0
mc_8x9	5	47.4	47.7	0.408	2/2/1

Every audited realization reproduces the frozen branch label of the original run; combined with the 0/185 discrepancies of the threshold re-classification of the reference campaign (Sec. S3), the branch taxonomy is internally consistent.

11 S11. Energy-matched comparison, trajectory tests, and B2 windows

Energy-matched LBM at $E^* = 0.25$, $Re = 500$ (integrated dissipation ε and ratio to the uniform reference):

Table 7: Iso-energetic dissipation ratios at $N = 256$ and $N = 512$.

Grid	$\varepsilon_{\text{unif}}$	$\varepsilon_{k=1}$	$\varepsilon_{k=2}$	ratios $k1/k2$
$N = 256$	9.4980×10^{-3}	6.8912×10^{-3}	1.1285×10^{-2}	0.7255 / 1.188
$N = 512$	1.10981×10^{-2}	7.1725×10^{-3}	1.18987×10^{-2}	0.6463 / 1.0721

The energy-matched ranking $k = 1 < \text{uniform} < k = 2$ is unchanged across resolutions; the $k = 1$ ratio range is $[0.646, 0.726]$ and the $k = 2$ (B_1) ratio range is $[1.072, 1.188]$.

Exp D trajectory tests ($Re = 500$, $E^* = 0.25$, probe spacing 50 epochs; seeds 4/7/8 $\leftrightarrow$ classes $B_3/B_1/B_2$). Control-energy autocorrelation $R_E(\Delta)$, control-cycle displacements D_1, D_4 , and modal displacements Δf :

Table 8: Exp D trajectory diagnostics.

seed	$R_E(1)$	$R_E(2)$	$R_E(4)$	D_1	D_4/D_1	$\Delta f(4)$ steady
4	+0.245	+0.167	+0.364	0.0744	2.92	0.0500
7	-0.446	+0.488	+0.372	0.0091	0.92	0.0011
8	+0.552	+0.821	+0.643	0.4178	2.29	0.0010

The four-epoch spectral peak of the raw per-epoch loss ($f \approx 0.25$; power up to $\sim 55\times$ the spectral median for seed 4) does not propagate to energy, control, or modal fractions: no period-4 cycle is detected in the control path, and branch labels stay stable across transients (seed 7: case A; seed 8: finite relocation transient within one class, case B; seed 4: slow persistent drift, case C). With $n = 60$ trajectory points ($n = 30$ stationary), the autocorrelation standard error is $\approx 0.13\text{--}0.18$.

B2 long-window quasi-stationarity ($Re = 500$, $E^* = 0.25$, $N = 256$; LBM run of 640 000 iterations, actuation period 256):

Table 9: B2 kinetic-energy fluctuation ratio over T and $2T$ windows.

window	cycles	$R_K = K_{\text{fluct}}/K_{\text{mean}}$	$ \Delta R_K /R_K(2T)$
T	128	14.95941409258121	—
$2T$	256	14.959414094980596	1.60×10^{-10}

with a 15% persistence threshold, the verdict is **PERSISTANT**: the B_2 flow is statistically stationary in the mean-dissipation sense while being strongly time-dependent at the actuation frequency ($R_K \simeq 15$).