

Supplementary Online Material for: Conditional modal branch selection in physics-informed active control of lid-driven cavity flow

Beniaiche Ahmed, Larabi Abderrahim, and Belkadi Mbarek
Ecole Militaire Polytechnique, Algiers, Algeria

This document provides the implementation details, complete datasets, and additional analyses supporting the main manuscript. Section numbers S1–S11 correspond to the references “Supplementary Sec. S...” used in the text.

1 S1. PINN architecture and training details

The parametric network represents the flow-control pair for the full Reynolds interval $[100, 1000]$ through inputs $(x, y, t, \widetilde{Re})$, with $\widetilde{Re} = (Re - 550)/450 \in [-1, 1]$.

- **Flow subnetwork:** 6 hidden layers $\times$ 64 neurons, tanh activation, 21 379 trainable parameters; outputs (u, v, p) .
- **Control subnetwork:** 4 hidden layers $\times$ 32 neurons, 3 614 trainable parameters; outputs the coefficients $\{a_{mn}\}$.
- **Precision:** all arithmetic in double precision; residuals by automatic differentiation.
- **Training:** 300-epoch pre-training (flow only, control frozen at zero), then 3 000-epoch joint phase with Adam ($\eta_0 = 8 \times 10^{-4}$, decay factor 0.995 per 100 epochs). Collocation points: $N_{\text{phys}} = 4\,000$ interior points resampled uniformly every epoch; $N_{\text{probe}} = 20$ for the Reynolds-separation term.
- **Loss weights:** $\lambda_{\text{PDE}} = 1$, $\lambda_{\text{BC}} = 15$, $\lambda_{\text{diss}} = 0.05$, $\lambda_{\text{ctrl}} = 200$, $\lambda_{\text{var}} = 1$, $\lambda_{Re} = 2$.

2 S2. Modal-energy quadrature conventions

For the lid-expansion $U_{\text{lid}}(x, t, Re) = \sum_m \sum_n a_{mn} \sin(m\pi x/L_x) \cos(n\pi t)$, the modal energy accounting uses $\langle \cos^2(0) \rangle = 1$ and $\langle \cos^2(n\pi t) \rangle = \frac{1}{2}$ for $n \geq 1$:

$$\mathcal{E}_{m0} = \frac{1}{2} a_{m0}^2, \quad \mathcal{E}_{mn} = \frac{1}{4} a_{mn}^2 \quad (n \geq 1).$$

The realized actuation energy is $E_r = \langle U_{\text{lid}}^2 \rangle_{s,t}$, with $s = x/L_x$, and the normalized modal fractions $f_{mn} = \mathcal{E}_{mn} / \sum \mathcal{E}_{ij}$ sum to unity. The stationary second-harmonic fraction is $f_2 \equiv f_{20}$ and the total time-dependent fraction is $f_t = \sum_m \sum_{n \geq 1} f_{mn}$.

Table 1: Pooled and per-condition statistics of the 185-realization campaign (30 cells; oversampled cells: $Re = 1000, E^* = 1$ ($n = 20$), $Re = 500, E^* = 0.25$ ($n = 15$), $Re = 700, E^* = 2$ ($n = 15$)).

Quantity	B_1	B_2	B_3
Pooled frequency (frozen labels)	42.2%	43.8%	14.1%
Pooled frequency (threshold re-classification)	42.2%	43.8%	14.1%
Mean per condition	42.1%	46.7%	11.3%
Pooled excluding oversampled cells ($N = 135$)	42.2%	48.1%	9.6%
Discrepancies frozen vs. threshold ($N = 185$)	0 / 185		

3 S3. Reference (Re, E^*) campaign – realization-level data

The threshold classification rule is: B_1 if $f_2 > 2f_t$ and $f_2 > 0.5$; B_2 if $f_t > 2f_2$ and $f_t > 0.5$; otherwise B_3 . The frozen labels and the threshold-based re-classification coincide for all 185 realizations (0/185). The realization-level file (Supplementary Table S1 in the submission system) lists, for each realization: (Re, E^*), seed, f_2 , f_t , the dominant coefficients, total loss, and frozen branch label.

4 S4. Variance-regularization campaign and Wilson intervals

Six values of $\lambda_{\text{var}} \in \{0, 0.25, 0.5, 1, 2, 5\}$ at three operating points, $N_s = 5$ seeds each (90 runs). Frequencies are reported in Table 3 of the main text. Finite-sample 95% Wilson-score confidence intervals use

$$p \pm \frac{z}{\sqrt{N_s}} \sqrt{\frac{\hat{p}(1 - \hat{p})}{N_s} + \frac{z^2}{4N_s^2}} / \left(1 + \frac{z^2}{N_s}\right), \quad z = 1.96.$$

For $N_s = 5$ and $\hat{p} = 0, 0.2, 0.4, 0.6, 1.0$ the half-widths are 0.31, 0.43, 0.47, 0.43, 0.31, respectively.

5 S5. Energy sweep and seed multiplicity

Energy sweep at $Re = 500$ ($E^* \in \{0.01, 0.05, 0.1, 0.25, 0.5, 1, 2\}$):

Table 2: Energy sweep at $Re = 500$.

E^*	0.01	0.05	0.1	0.25	0.5	1	2
f_2 (%)	~ 94	~ 93	92.2	92.2	–	–	74.0
f_t (%)	~ 0.5	~ 0.7	1.1	1.3	–	–	21.8

Seed multiplicity at $Re = 500$, $E^* = 0.25$, 10 seeds: 4 B_1 , 4 B_2 , 2 B_3 ; total losses in [1.371, 1.394] (relative spread 1.7%).

6 S6. Capacity and truncation experiments

Capacity variants at $Re = 500$, $E^* = 0.25$: 4×32 (small), 6×64 (reference), 8×128 (large) give $f_2 \in [89.5\%, 92.2\%]$. Control truncations $(N_x, N_t) \in \{(4, 3), (6, 1), (6, 3), (6, 5), (8, 5), (8, 9)\}$ give $f_2 \in [87.3\%, 92.3\%]$ with $A_2 \in [0.668, 0.684]$; see Table 6 of the main text.

7 S7. Spectral bases and aspect ratio

Fourier basis: sine in space (enforces corner compatibility) and cosine in time ($n = 0 = \text{stationary}$). Modified-Chebyshev basis: same structure with Chebyshev polynomials of the first kind T_m on the lid interval, normalized to the same energy budget. Aspect-ratio cells $AR \in \{0.5, 1, 2\}$ at $Re = 500$, $E^* = 0.25$: $f_2 = 69.9\%$, 92.2% , 91.2% (see Table 7 of the main text).

8 S8. MRT-LBM solver and grid-convergence index (GCI)

Solver. D2Q9 lattice with the diagonal relaxation vector $\mathbf{S} = \text{diag}(0, 1.4, 1.4, 0, 1.2, 0, 1.2, \omega, \omega)$, $\omega = 1/\tau$, $\tau = 3\nu_{\text{LBM}} + 0.5$, $\nu_{\text{LBM}} = U_{\text{ref}}N/Re$ with $U_{\text{ref}} = 0.05$. Lid forcing through a velocity boundary condition; boundary populations via the Zou–He scheme. Parallelism via Numba **prange** over lattice rows.

Benchmark. Centerline velocity errors against Ghia et al. at $N = 256$: $L_2(u) = 0.00506$, $L_\infty(u) = 0.01283$ ($Re = 100$); $L_2(u) = 0.01411$, $L_\infty(u) = 0.03003$ ($Re = 1000$); corresponding v -component errors $0.00675/0.01136$ and $0.01321/0.02125$.

GCI. Three-grid study of the $A_2 \sin(2\pi x)$ control at $Re = 500$, $N \in \{128, 256, 512\}$:

$$\varepsilon_{128} = 0.0093500, \quad \varepsilon_{256} = 0.0104070, \quad \varepsilon_{512} = 0.0109658.$$

Observed order $p_{\text{obs}} = 0.9196$ (from $p = \ln |(\varepsilon_{256} - \varepsilon_{128})/(\varepsilon_{512} - \varepsilon_{256})| / \ln 2$), Richardson extrapolation $\varepsilon_\infty = 0.01159$. The GCI (formula S8.4) uses the safety factor $F_s = 1.25$ and refinement ratio $r = 2$:

$$\text{GCI}_j = F_s \frac{|\varepsilon_{j+1} - \varepsilon_j|}{(r^p - 1) \varepsilon_j}, \quad \Rightarrow \quad \text{GCI}_{512} = 7.15\%, \quad \text{GCI}_{256} = 14.25\%.$$

Uniform-to-extrapolated reduction: $0.01159/0.0442824 = 73.8\%$ (vs. 76.5% measured at $N = 256$).

9 S9. Symbolic regression: configuration and LOCO scores

PySR over binary operators $\{+, -, \times, /\}$ and unary $\{\sin, \cos, \exp, \text{square}\}$, 20 populations, 100 iterations each, Pareto selection on complexity–accuracy trade-off. Best expression: $U_{\text{SR}} \approx a[b + \cos(\omega_{\text{SR}}x - \phi)]^2(c_0 + c_1x + c_2t + c_3Re) + c_4$, $\omega_{\text{SR}} \approx 3.42$. Leave-one-Reynolds-condition-out R^2 over $Re \in \{100, 300, 400, 500, 600, 700, 800, 900, 1000\}$: minimum 0.9326 ($Re = 900$), maximum 0.9908 ($Re = 500$); all nine values above 0.93 .

10 S10. Reproducibility double runs and multi-seed audit

Double runs A/B (idempotent, same seeds/hyperparameters, base 6×5 , $E^* = 0.25$, final code) at $Re = 500$:

Runs A and B are bit-identical to the last decimal for seeds 7 and 8, and quasi-identical for seed 4 (same class B_3). All three coincide with the original realizations re-analyzed with the final code.

Multi-seed audit (99 realizations, final code, 18 configurations):

Every audited realization reproduces the frozen branch label of the original run; combined with the $0/185$ discrepancies of the threshold re-classification (Sec. S3), the branch taxonomy is internally consistent.

Table 3: Reproducibility double-run results (A/B/original).

seed	run	E	f_2	f_t	A_{10}	A_{20}	branch
4	A	0.2514	0.5957	0.3693	0.0376	-0.5472	B_3
4	B	0.2512	0.6080	0.3575	0.0394	-0.5527	B_3
4	orig	0.2495	0.5215	0.4501	0.0364	-0.5101	B_3
7	A	0.2508	0.9271	0.0083	-0.0604	0.6819	B_1
7	B	0.2508	0.9271	0.0083	-0.0604	0.6819	B_1
7	orig	0.2498	0.9131	0.0162	-0.0681	0.6754	B_1
8	A	0.2507	0.0357	0.9609	-0.0109	0.1339	B_2
8	B	0.2507	0.0357	0.9609	-0.0109	0.1339	B_2
8	orig	0.2503	0.0133	0.9841	-0.0158	0.0817	B_2

Table 4: Audit summary per configuration: mean f_2 , mean f_t , mean A_2 (in %), and branch counts.

label	n	$\bar{f}_2$	$\bar{f}_t$	$\bar{A}_2$	$B_1/B_2/B_3$
ar_rect_2	5	90.8	1.2	0.687	5/0/0
ar_rect_half	5	15.3	83.8	0.187	1/4/0
ar_square	5	73.0	20.1	0.572	4/1/0
bas_cheb_re100	5	0.5	99.2	-0.030	0/5/0
bas_cheb_re1000	5	1.2	98.4	-0.047	0/5/0
bas_cheb_re500	5	2.1	97.6	-0.085	0/5/0
bas_fourier_re100	5	32.6	64.7	0.313	1/3/1
bas_fourier_re1000	5	73.5	21.3	0.610	3/0/2
bas_fourier_re500	5	73.0	20.1	0.572	4/1/0
cap_baseline	8	64.0	30.5	0.402	5/2/1
cap_large	8	89.1	3.6	0.676	8/0/0
cap_small	8	92.2	1.3	0.694	8/0/0
mc_4x3	5	54.4	39.4	0.436	3/2/0
mc_6x1	5	91.1	0.0	0.683	5/0/0
mc_6x3	5	75.1	16.1	0.485	4/1/0
mc_6x5	5	73.1	20.1	0.572	4/1/0
mc_8x5	5	59.1	34.9	0.485	3/2/0
mc_8x9	5	47.4	47.7	0.408	2/2/1

11 S11. Energy-matched comparison, trajectory tests, and B2 windows

Energy-matched LBM at $E^* = 0.25$, $Re = 500$ (integrated dissipation ε and ratio to the uniform reference):

Table 5: Iso-energetic dissipation ratios at $N = 256$ and $N = 512$.

Grid	$\varepsilon_{\text{unif}}$	$\varepsilon_{k=1}$	$\varepsilon_{k=2}$	ratios $k1/k2$
$N = 256$	9.4980×10^{-3}	6.8912×10^{-3}	1.1285×10^{-2}	0.7255 / 1.188
$N = 512$	1.10981×10^{-2}	7.1725×10^{-3}	1.18987×10^{-2}	0.6463 / 1.0721

The energy-matched ranking $k = 1 < \text{uniform} < k = 2$ is unchanged across resolutions; the $k = 1$ ratio range is $[0.646, 0.726]$ and the $k = 2$ (B_1) ratio range is $[1.072, 1.188]$.

Exp D trajectory tests ($Re = 500$, $E^* = 0.25$, probe spacing 50 epochs; seeds 4/7/8 $\leftrightarrow$ classes $B_3/B_1/B_2$). Control-energy autocorrelation $R_E(\Delta)$, control-cycle displacements D_1, D_4 ,

and modal displacements Δf :

Table 6: Exp D trajectory diagnostics.

seed	$R_E(1)$	$R_E(2)$	$R_E(4)$	D_1	D_4/D_1	$\Delta f(4)$ steady
4	+0.245	+0.167	+0.364	0.0744	2.92	0.0500
7	−0.446	+0.488	+0.372	0.0091	0.92	0.0011
8	+0.552	+0.821	+0.643	0.4178	2.29	0.0010

The four-epoch spectral peak of the raw per-epoch loss ($f \approx 0.25$; power up to $\sim 55\times$ the spectral median for seed 4) does not propagate to energy, control, or modal fractions: no period-4 cycle is detected in the control path, and branch labels stay stable across transients (seed 7: case A; seed 8: finite relocation transient within one class, case B; seed 4: slow persistent drift, case C). With $n = 60$ trajectory points ($n = 30$ stationary), the autocorrelation standard error is ≈ 0.13 – 0.18 .

B2 long-window quasi-stationarity ($Re = 500$, $E^* = 0.25$, $N = 256$; LBM run of 640 000 iterations, actuation period 256):

Table 7: B2 kinetic-energy fluctuation ratio over T and $2T$ windows.

window	cycles	$R_K = K_{\text{fluct}}/K_{\text{mean}}$	$ \Delta R_K /R_K(2T)$
T	128	14.95941409258121	—
$2T$	256	14.959414094980596	1.60×10^{-10}

with a 15% persistence threshold, the verdict is **PERSISTANT**: the B_2 flow is statistically stationary in the mean-dissipation sense while being strongly time-dependent at the actuation frequency ($R_K \simeq 15$).